

# Metaheuristics: Single State Local Search Methods

Fernando Lobo

Universidade do Algarve

# Stochastic Local Search

- Almost all metaheuristics make use of randomized choices
- This gives rise to a family of algorithms known as *Stochastic Local Search* (SLS) methods
- Stochastic is just a fancy name for randomized
- Note: making randomized choices doesn't mean the algorithm is blind and searches completely at random

# Stochastic Local Search

- SLS methods often use of the notion of *locality*
- Locality is induced by the neighbourhood of solutions
- That is, SLS tend to search around the vicinity of solution(s) that the algorithm maintain(s) as its current state
  - ▶ Single-state SLS methods only maintain one solution as its current state
  - ▶ Multi-state SLS methods maintain several solutions as its current state (they are also called *population-based search methods*)
- We shall study both types of algorithms

# Single-state SLS methods

Let's start with the most basic methods:

- Blind search
  - ▶ Uninformed Random Walk
  - ▶ Uninformed Random Picking
- Hillclimbing
  - ▶ Steepest Ascent (Descent)
  - ▶ Next Ascent (Descent)

# Blind Search

- Starting from an initial solution, the algorithm wanders around the search space and simply keeps the best solution found so far
- Two variations:
  - ▶ Uninformed Random Walk
    - ★ At each step the algorithm moves from a solution  $s$  to a solution in the neighbourhood of  $s$
  - ▶ Uninformed Random Picking
    - ★ At each step the algorithm moves from a solution  $s$  to any other solution in the search space
- Performs poorly

# Uninformed Random Walk

---

$s$  = solution generated uniformly at random (u.a.r.)

$v_s = f(s)$

$best = (s, v_s)$

**while** *not terminate()* **do**

$x$  = choose solution u.a.r. from  $\mathcal{N}(s)$

$v_x = f(x)$

**if**  $v_x > v_s$  **then**

$best = (x, v_x)$

**end**

$s = x$

**end**

**return**  $best$

---

# Uninformed Random Picking

---

$s$  = solution generated uniformly at random (u.a.r.)

$v_s = f(s)$

$best = (s, v_s)$

**while** *not terminate()* **do**

$x$  = choose solution u.a.r. from **S**

$v_x = f(x)$

**if**  $v_x > v_s$  **then**

$best = (x, v_x)$

**end**

$s = x$

**end**

**return**  $best$

---

# Some notes on the pseudocode

- u.a.r. stands for uniformly at random
- $f$  denotes the evaluation function
- `terminate()` is a boolean predicate
  - ▶ could be based on an allowed maximum number of iterations
  - ▶ a certain amount of time
  - ▶ or some other criterion
- algorithm assumes we are dealing with a maximization problem
  - ▶ for minimization, change  $>$  to  $<$



# Hillclimbing

- Maintain a current solution
- At each time step, move to a neighbour that is better than the current solution
- If no such neighbour exist, terminate
  - ▶ we have reached a local optimum

# Hillclimbing

- Two common variations:
  - ▶ *Steepest Ascent*
    - ★ visit the entire neighbourhood and move to the best neighbour (with ties broken u.a.r.)
  - ▶ *Next Ascent*
    - ★ visit the neighbourhood in a random order, and move to the first neighbour that improves upon the current solution
- Steepest Ascent requires the evaluation of the entire neighbourhood in every search step
- Because of that, Next Ascent is often much more efficient than Steepest Ascent

# Steepest Ascent Hillclimbing

---

$s$  = solution generated u.a.r.

$v_s = f(s)$

**repeat**

$best\_improvement = 0$

**foreach**  $x \in \mathcal{N}(s)$  **do**

$v_x = f(x)$

**if**  $v_x - v_s > best\_improvement$  **then**

$best\_improvement = v_x - v_s$

$best\_neighbour = (x, v_x)$

**end**

**end**

**if**  $best\_improvement > 0$  **then**

$(s, v_s) = best\_neighbour$

**end**

**until**  $best\_improvement = 0$

**return**  $(s, v_s)$

---

# Next Ascent Hillclimbing

---

$s$  = solution generated u.a.r.

$v_s = f(s)$

**repeat**

$improve = \text{false}$

**foreach**  $x \in \mathcal{N}(s)$  **do**

$v_x = f(x)$

**if**  $v_x > v_s$  **then**

$improve = \text{true}$

$(s, v_s) = (x, v_x)$

**break**

**end**


**end**

**until** *not improve*

**return**  $(s, v_s)$

---

## A 1-dimensional multimodal function



# A 2-dimensional multimodal function




Figure taken from the GECCO 2020 Competition on Niching Methods for Multimodal Optimization website located at <https://www.epitropakis.co.uk/gecco2020/>

# Multi-start Hillclimbing

- Drawback of hillclimbing: getting stuck at a local optimum
- Quick fix: once a local optimum is reached, restart hillclimbing again
- Hopefully we will reach another local optimum
- We can iterate this over and over, obtaining a collection of local optima
- The resulting algorithm is called *Multi-start Hillclimbing*

# Another strategy to escape a local optimum

- Use larger neighbourhoods
- SAT example: 2-bit Hamming distance neighbourhood, instead of 1-bit
- The resulting local optimum would be better than any of its 2-bit neighbours
- But this has a cost:
  - ▶ time to explore a larger neighbourhood increases
  - ▶ in theory, we could make the neighbourhood so large that its size becomes of the order of the size of the search space!
  - ▶ need to find a balance between time complexity and ability to escape local optima



# Neighbourhood pruning

- A strategy to improve efficiency
- Reduce size of neighbourhood by excluding neighbours that cannot (or are unlikely to) result in an improvement

# Neighbourhood pruning: TSP example under 2-exchange

- 2 edges are removed, and 2 edges are added
- To be an improving move, at least one of the added edges has to have a smaller distance (weight) than a removed edge
  - ▶ Why? Because the contribution of the other edges are held the same in both solutions
- We can use this knowledge and avoid visiting (guaranteed) non-improving neighbours
- We can even build an appropriate data structure to support this

# Neighbourhood pruning: TSP example under 2-exchange (candidate lists)

## Candidate Lists:

- For each vertex store a list of neighbouring vertices sorted according to edge weight
- These are called *candidate lists*
- Building complete candidate lists requires  $O(n^2)$  memory and  $O(n^2 \log n)$  time
- To avoid this, the size of the candidate lists is often bounded by a constant value: 10 to 40

# Variable Neighbourhood Ascent

- Idea: when stuck at a local optimum, switch to a larger neighbourhood
- Use  $k$  neighbourhoods  $\mathcal{N}_1, \mathcal{N}_2, \dots, \mathcal{N}_k$  in increasing order of size
- When an improvement is made using neighbourhood  $\mathcal{N}_j$ , we go back to using  $\mathcal{N}_1$
- At the end, the resulting solution will be a local optimum with respect to all these neighbourhoods